



Kinect for Windows API V2

John Elsbree

Principal Software Development Engineer
Kinect for Windows Program

18 November 2013



Microsoft Confidential



Everything is under NDA

Unless otherwise stated.



Agenda

- V2 Kinect sensor
- API overview
 - Influences and style
 - Differences from V1
- API features
 - Sensor lifecycle
 - Data sources
 - Frame synchronization
 - Coordinate mapping
- Porting guidance
- Not covered in this session: speech, face tracking, interactions



V2 Kinect sensor

Sensor differences from V1 (briefly)

- One color camera resolution (1920x1080), frame rate (30 Hz)
- One depth/IR camera resolution (512x424), frame rate (30 Hz)
- Depth range from 0.5 to 4.5 m
- Clean infrared frames
- Can use infrared and color cameras simultaneously
- No tilt motor (wider field of view)




API overview

COM and .NET

- V2 API still has both COM and .NET flavors
- Much more similar *to each other* than they were in V1
- COM and .NET each have unique patterns for expressing:
 - Events
 - Buffers
 - Collections
- V1 code *cannot* just be recompiled for V2; some changes will be needed
- Existing .NET code will need fewer changes than COM code

Influences and style

- New API is influenced by WinRT
- Xbox One has a WinRT API for Kinect
- Kinect for Windows V2 API is **not** a WinRT API (yet)
- “Shape” of V2 COM and .NET APIs identical to Xbox One’s WinRT API
 - Same types, methods, etc.
- New Kinect API uses WinRT idioms
 - Sources/readers
 - Events
 - Collections




API differences from V1

- All resolutions and frame rates are constant
- Can use color and infrared simultaneously
- No separate "near" mode for depth (range is 0.5 to 4.5 m)
- "Skeleton" → "Body"
 - 6 fully-tracked bodies
 - More joints per body
 - More features: expressions, actions, lean, ...
 - No separate "seated" mode (both seated and standing can be tracked)
- Depth and body index delivered as separate frame types
- More audio beam features

Source/Reader pattern

- V1: Once a frame was retrieved from a stream by polling, it was gone forever
- V2: Multiple reader instances can independently poll the same source for frames, without interference
- Readers can be paused/resumed independently
- Enables more componentized applications



Source/Reader pattern

- V1: sensor → stream → frame → data
 - Stream – 1 of each type per sensor
- V2: sensor → source → reader → frame → data
 - Source – 1 of each type per sensor
 - Reader – many per source

Events



- COM:
 - HRESULT *SubscribeEventName*(WAITABLE_HANDLE *waitableHandle);
 - HRESULT *UnsubscribeEventName*(WAITABLE_HANDLE waitableHandle);
 - HRESULT *GetEventNameEventData*(WAITABLE_HANDLE waitableHandle, *IEventNameEventArgs* **eventData);
- .NET:
 - public event EventHandler<*EventNameEventArgs*> *EventName*;

Buffers

- COM:
 - HRESULT CopyFrameDataToArray(UINT capacity, BYTE* buffer);
- .NET:
 - public void CopyFrameDataToArray(byte[] frameData);
 - public void CopyFrameDataToBuffer(uint bufferSize, IntPtr buffer);



Collections

- COM:
 - `interface IKinectSensorCollection`
- .NET:
 - `class KinectSensorCollection :
 IReadOnlyList<KinectSensor>,
 INotifyCollectionChanged`



API Features

Sensor lifecycle

- Lifecycle
 - Find a KinectSensor object
 - Open it
 - Use it
 - Close it
- Sensor unplugged:
 - KinectSensor object remains valid
 - Your code still runs
 - No frames arrive
 - KinectSensor.IsAvailable tells you if it's actually there



Sensor lifecycle – One sensor



```
this.sensor = KinectSensor.Default;  
this.sensor.Open();  
...  
this.sensor.Close();
```

Sensor lifecycle – Multiple sensors

- **NOTE:** Not yet implemented in Tech Preview

```
foreach (KinectSensor sensor in KinectSensor.Sensors)
{
    string sensorId = sensor.UniqueKinectId;
    this.sensors.Add(sensorId, sensor);


    sensor.Open();
    ...
}
...

foreach (KinectSensor sensor in this.sensors.Values)
{
    sensor.Close();
}
```



Data sources

- Infrared
- Color
- Depth
- Body index
- Body
- Audio



Infrared

- The simplest sources to use
- Two forms: Infrared and LongExposureInfrared
 - Infrared is a single frame
 - LongExposureInfrared is sum of 3 IR frames (higher signal:noise ratio, but more motion blur)
- Examples that follow use single-frame Infrared
 - To use long-exposure infrared, just "Infrared" → "LongExposureInfrared"
- Frame data is 2 bytes per pixel: 16-bit IR intensity value (same as V1)



Infrared – Initialization (.NET)

```
// Allocate a buffer
FrameDescription frameDesc =
    this.sensor.InfraredFrameSource.FrameDescription;
this.infraredData = new ushort[frameDesc.LengthInPixels];

// Open a reader and subscribe to frame arrival events
this.infraredReader = this.sensor.InfraredFrameSource.OpenReader();
this.infraredReader.FrameArrived += InfraredFrameArrived;
```

Aside – FrameDescription

```
public sealed class FrameDescription
{
    public int Width { get; }
    public int Height { get; }
    public float HorizontalFieldOfView { get; }
    public float VerticalFieldOfView { get; }
    public float DiagonalFieldOfView { get; }
    public uint LengthInPixels { get; }
    public uint BytesPerPixel { get; }
}
```

Infrared – Event handler (.NET)

```
private void InfraredFrameArrived(
    object sender,
    InfraredFrameArrivedEventArgs e)
{
    bool processFrame = false;

    // Acquire the frame
    using (InfraredFrame frame = e.FrameReference.AcquireFrame())
    {
        if (null != frame)
        {
            // Copy frame's data to our buffer
            frame.CopyFrameDataToArray(this.infraredData);
            processFrame = true;
        }
    }

    if (processFrame) { ProcessInfraredData(this.infraredData); }
}
```

Aside – FrameReference

- `FrameReference.RelativeTime` property is a timestamp for the frame it represents
- By the time the app calls `FrameReference.AcquireFrame`, the frame might already have “expired”
 - A newer frame has already taken its place
 - `AcquireFrame` returns null
 - Caller always needs to check

Infrared – Initialization (COM)

```
// Access the frame source
IIInfraredFrameSource* pSource = nullptr;
this->_pSensor->get_InfraredFrameSource(&pSource);

// Allocate a buffer
IFrameDescription* pFrameDesc = nullptr;
pSource->get_FrameDescription(&pFrameDesc);
pFrameDesc->get_LengthInPixels(&this->_lengthInPixels);
pFrameDesc->Release();
this->_pInfraredData = new UINT16[this->_lengthInPixels];

// Open a reader and subscribe to frame arrival events
pSource->OpenReader(&this->_pInfraredReader);
this->_pInfraredReader->SubscribeFrameArrived(&this->_hInfraredEvent);
pSource->Release();
```

Infrared – Event dispatch (COM)

```
while (...)
{
    HANDLE handles[] = { reinterpret_cast<HANDLE>(this->_hInfraredEvent), ... };
    switch (WaitForMultipleObjects(_countof(handles), handles, ...))
    {
        case WAIT_OBJECT_0:
        {
            IInfraredFrameArrivedEventArgs* pArgs = nullptr;
            this->_pInfraredReader->GetFrameArrivedEventData(this->_hInfraredEvent,
                &pArgs);
            InfraredFrameArrived(pArgs);
            pArgs->Release();
        }
        break;
        ...
    }
}
```

Infrared – Event handler (COM)

```
void MyClass::InfraredFrameArrived(IInfraredFrameArrivedEventArgs* pArgs)
{
    IInfraredFrameReference* pFrameReference = nullptr;
    pArgs->get_FrameReference(&pFrameReference);

    // Acquire the frame
    bool processFrame = false;
    IInfraredFrame* pFrame = nullptr;
    if (SUCCEEDED(pFrameReference->AcquireFrame(&pFrame)))
    {
        // Copy frame's data to our buffer
        pFrame->CopyFrameDataToArray(this->_lengthInPixels, this->_pInfraredData);
        processFrame = true;
        pFrame->Release();
    }
    pFrameReference->Release();

    if (processFrame) { ProcessInfraredData(this->_lengthInPixels, this->_pInfraredData); }
}
```

Infrared – Other features (.NET)

- Polling (instead of events)

```
InfraredFrame frame =  
    this.infraredReader.AcquireLatestFrame();
```

- Pause/resume

```
this.infraredReader.IsPaused = true;
```

- Raw buffer access

```
int size;  
IntPtr buffer;  
frame.AccessUnderlyingBuffer(out size, out buffer);  
unsafe { ushort* bufferData = (ushort*)buffer; ... }
```

- Timestamp

```
long timestamp = frame.RelativeTime;
```

Color

- Color frames have multiple possible formats (RGBA, BGRA, YUY2, ...)
- Frames come from the sensor in a default raw format
 - YUY2 for now (Tech Preview), but may be different in the future
- Frame data can be:
 - Accessed in its raw format, or
 - Converted to another format (at slightly higher cost)
- Underlying buffer access is available only for the raw format
- Buffer is array of bytes, but typically multiple bytes per pixel (how many depends on format)



Color – Raw format

```
FrameDescription frameDesc =
    this.sensor.ColorFrameSource.FrameDescription;
this.colorData =
    new byte[frameDesc.LengthInPixels * frameDesc.BytesPerPixel];

ColorImageFormat rawColorFormat = ColorImageFormat.None;
using (ColorFrame frame = e.FrameReference.AcquireFrame())
{
    if (null != frame)
    {
        rawColorFormat = frame.RawColorImageFormat;
        frame.CopyRawFrameDataToArray(this.colorData);
    }
}
switch (rawColorFormat) ...
```

Color – Format conversion

```
FrameDescription frameDesc =  
    this.sensor.ColorSource.CreateFrameDescription(ColorImageFormat.Bgra);  
this.colorData =  
    new byte[frameDesc.LengthInPixels * frameDesc.BytesPerPixel];
```

```
using (ColorFrame frame = e.FrameReference.AcquireFrame())  
{  
    if (null != frame)  
    {  
        frame.CopyConvertedFrameDataToArray(  
            this.colorData,  
            ColorImageFormat.Bgra);  
    }  
}
```

Color – Convert directly to WriteableBitmap

```
// Determine dimensions from frame description
FrameDescription frameDesc = colorFrame.CreateFrameDescription(ColorImageFormat.Bgra);
int width = frameDesc.Width;
int height = frameDesc.Height;
uint bufferLength = (uint)(frameDesc.LengthInPixels * frameDesc.BytesPerPixel);

// Create bitmap if needed
if (null == this.colorBitmap)
{
    this.colorBitmap = new WriteableBitmap(width, height, 96, 96, PixelFormats.Bgra32, null);
}

// Convert color data directly into the bitmap
this.colorBitmap.Lock();
colorFrame.CopyConvertedFrameDataToBuffer(
    bufferLength, this.colorBitmap.BackBuffer, ColorImageFormat.Bgra);
this.colorBitmap.AddDirtyRect(new Int32Rect(0, 0, width, height));
this.colorBitmap.Unlock();
```


Depth

- Frame data is 2 bytes per pixel: 16-bit distance in millimeters
- No “PlayerIndex” – use BodyIndex source instead
- API nearly identical to Infrared
- Two additional properties:
 - DepthMinReliableDistance
 - DepthMaxReliableDistance
- Known bug in Tech Preview:
DepthMaxReliableDistance is 4000,
but should be 4500 (depth values up
to 4500 mm are actually returned)




Body index

- Frame data is 1 byte per pixel: index of the body, as determined by body tracking
- Resolution is same as depth
- Pixel values (different from V1)
 - -1: No body at this pixel
 - 0 to 5: Index of the corresponding body, as tracked by the body source
 - All other values: not used
- Apart from pixel value size, API nearly identical to Infrared



Body

- Formerly known (in V1) as Skeleton
- Frame data is an array of Body objects
- Many new features in V2
 - More joints (neck, thumbs, hand tips)
 - Hand states (open, closed, "lasso")
 - Activities (eye closed, mouth open, mouth motion, looking away)
 - Appearance (wearing glasses)
 - Level of user engagement
 - Facial expressions (happy, neutral)
 - Lean direction (2D vector, "human joystick")



Body - Initialization

- Instead of FrameDescription, we have BodyCount

```
// Allocate a buffer of bodies
this.infraredData =
    new Body[this.sensor.BodyFrameSource.BodyCount];

// Open a reader and subscribe to frame arrival events
this.bodyReader = this.sensor.BodySource.OpenReader();
this.bodyReader.FrameArrived += BodyFrameArrived;
```

Body – Event handler

- Instead of CopyFrameDataToArray, we have GetAndRefreshBodyData


```
// Acquire the frame
using (BodyFrame frame = e.FrameReference.AcquireFrame())
{
    if (null != frame)
    {
        // Copy frame's data to our buffer
        frame.GetAndRefreshBodyData(this.bodyData);
    }
}
```

- Frame also has BodyCount and FloorClipPlane properties

Body – GetAndRefreshBodyData


- Designed to minimize per-frame allocations of new objects
- Your responsibility: pass a Body array of the appropriate length (BodyCount)
- If an array element is null:
 - New Body is created and populated with data, and
 - Stored in the array
- If an array element is non-null:
 - Existing Body instance is reused, and
 - Content of the Body is overwritten with data for the new frame
- To retain a Body after its frame is disposed:
 - Keep a reference to it, and
 - Replace it in the array with null

Body - Joints



Body – Joints positions and orientations

- Two dictionaries, each keyed by JointType:
 - Joints: tracking state and 3D position
 - Orientation: 3D orientation, specified as a quaternion
- Joint tracking state may be:
 - Not tracked
 - Inferred
 - Tracked



Body – Joint positions and orientations

```
foreach (Body body in this.bodyData)
{
    if (body.IsTracked)
    {
        foreach (Joint joint in body.Joints.Values)
        {
            if (joint.TrackingState != TrackingState.NotTracked)
            {
                CameraSpacePoint point = joint.Position;
                Vector4 rotation = body.JointOrientations[joint.JointType];
                ...
            }
        }
    }
}
```

Body – Hand states

- Two properties: HandLeftState, HandRightState
 - Unknown
 - Not tracked
 - Open
 - Closed
 - Lasso
- Confidence properties: HandLeftConfidence, HandRightConfidence
 - High
 - Low
- Hand state tracking limited to 2 bodies at a time
 - BodyFrameSource.OverrideHandTracking lets you choose which bodies
 - Limit *may* be increased in the future



Body – Activities, Appearance, Expressions

- Activities
 - EyeLeftClosed
 - EyeRightClosed
 - MouthOpen
 - MouthMoved
 - LookingAway
- Appearance
 - WearingGlasses
- Expressions
 - Happy
 - Neutral



Body – Activities, Appearance, Expressions

- Three dictionaries
 - Key is a state type
 - Value indicates probability of that state (Unknown, No, Maybe, Yes)
- API may be extended in the future, by adding new state keys

```
bool isMouthClosed =  
    (body.Activities[Activity.MouthOpen] == DetectionResult.No);  
  
bool isWearingGlasses =  
    (body.Appearance[Appearance.WearingGlasses] == DetectionResult.Yes);  
  
bool isPossiblyHappy =  
    (body.Expressions[Expression.Happy] >= DetectionResult.Maybe);
```

Body – Other properties

- Engaged

```
bool isEngaged = (body.Engaged == DetectionResult.Yes);
```

- Lean (body as a 2D joystick)

```
if (body.LeanTrackingState == TrackingState.tracked)
{
    float leanLeftRight = body.Lean.X;
    float leanForwardBack = body.Lean.Y;
}
```

- TrackingId: unique 64-bit ID assigned to each new body
- ClippedEdges: which edges of the field-of-view are clipping the body

Audio

- **NOTE:** Not yet implemented in Tech Preview
- Audio beam
 - A steerable “cone” of focus for audio
 - May be automatically or manually aimed
 - Kinect audio source API can support multiple beams
- Audio beam frame
 - Contains audio samples captured for a beam over a specific interval of time
 - Synchronized frames for all beams available as a group, on each event



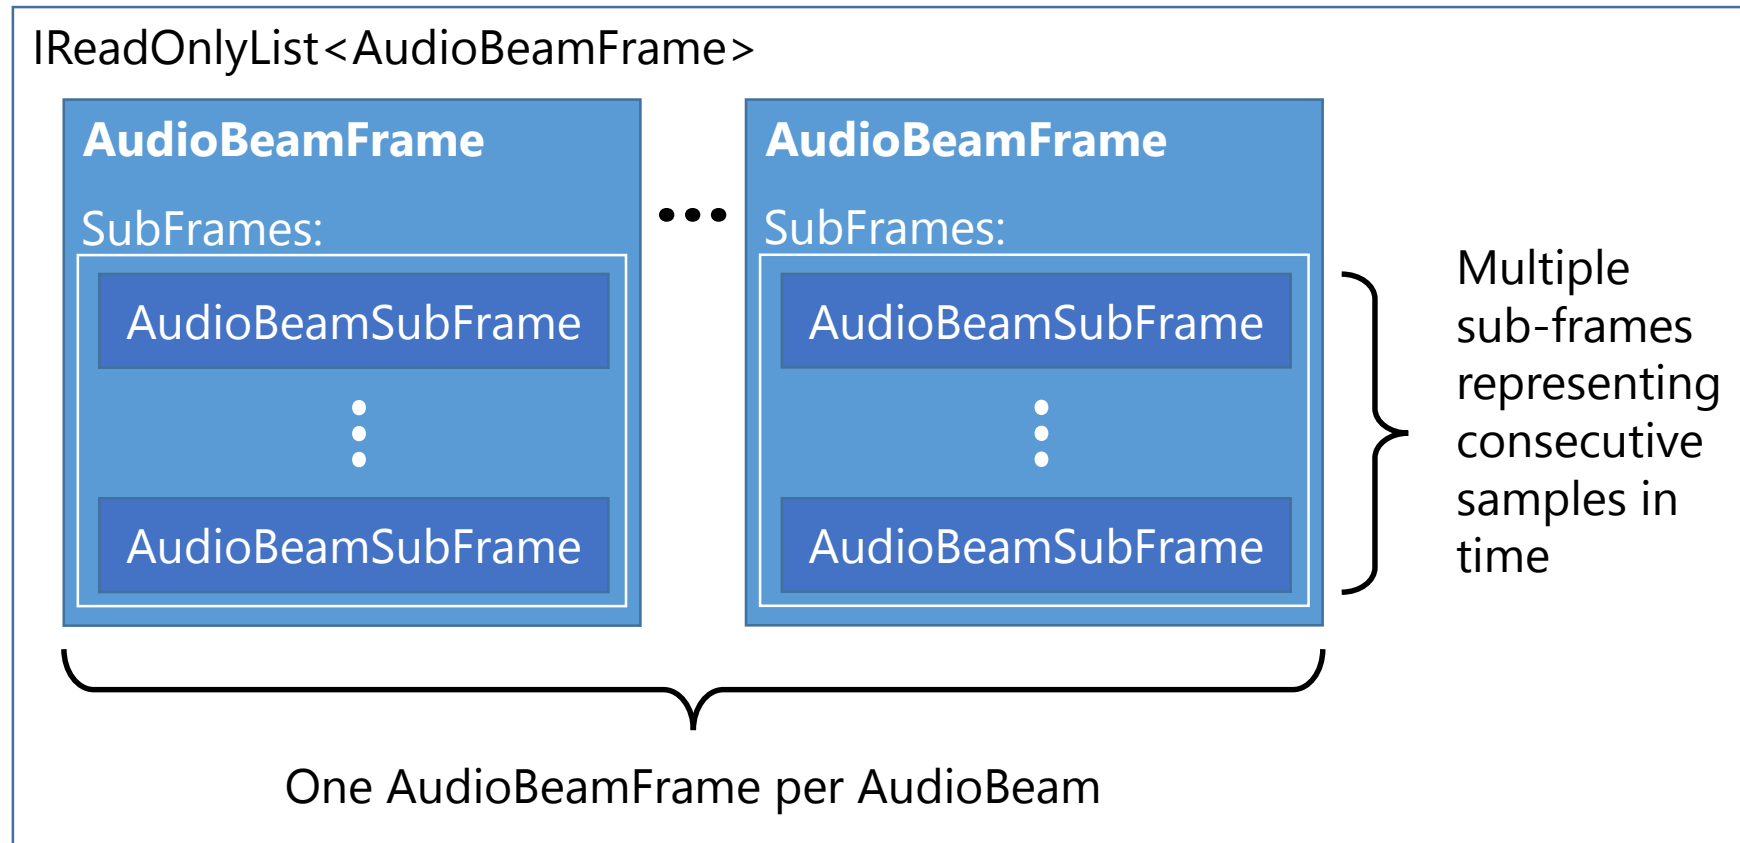
Audio – Single beam – Initialization

```
// Allocate a buffer
this.audioData = new
    byte[this.sensor.AudioSource.SubFrameLengthInBytes];

// Open a reader and subscribe to frame arrival events
this.audioReader = this.sensor.AudioSource.OpenReader();
this.audioReader.FrameArrived += AudioFrameArrived;
```

Audio – Beam frames

AudioBeamFrameReference.AcquireBeamFrames returns a list of AudioBeamFrames



Audio – Single beam – Event handler

```
void AudioFrameArrived(object sender, AudioBeamFrameArrivedEventArgs e)
{
    using (AudioBeamFrame beamFrame =
        e.FrameReference.AcquireBeamFrames().FirstOrDefault())
    {
        if (beamFrame != null)
        {
            foreach (AudioBeamSubFrame subFrame in beamFrame.SubFrames)
            {
                subFrame.CopyFrameDataToArray(this.audioData);
                float beamAngle = subFrame.BeamAngle;
                float beamAngleConfidence = subFrame.BeamAngleConfidence;
                long timestamp = subFrame.RelativeTime;
                subFrame.Dispose();

                ProcessAudioData(this.audioData, timestamp, beamAngle, beamAngleConfidence);
            }
        }
    }
}
```

Audio – Body correlation


- Identifies which body (or bodies) are in the path of a beam

```
foreach (AudioBeamSubFrame subFrame in beamFrame.SubFrames)
{
    subFrame.CopyFrameDataToArray(this.audioData);
    long timestamp = subFrame.RelativeTime;

    foreach (AudioBodyCorrelation audioBody in
        subFrame.AudioBodyCorrelations)
    {
        ProcessAudioData(
            this.audioData, timestamp, audioBody.BodyTrackingId);
    }

    subFrame.Dispose();
}
```

Audio – Beam steering



```
private int ManuallyAimAudioBeams(float[] targetAngles)
{
    int i = 0;
    foreach (AudioBeam beam in
        this.sensor.AudioSource.AudioBeams)
    {
        if (i >= targetAngles.Length) { break; }

        beam.AudioBeamMode = AudioBeamMode.Manual;
        beam.BeamAngle = targetAngles[i++];
    }

    return i; // number of beams that were actually aimed
}
```

Frame synchronization

- Open a MultiSourceFrameReader, indicating which sources you want
- When a matched set of frames is ready, the event fires
- MultiSourceFrame contains references to each of the matched frames
- A more general-purpose form of the AllFramesReady event in V1



MultiSourceFrameReader - Initialization

```
// Open a multi-source reader
this.multiReader =
    this.sensor.OpenMultiSourceFrameReader(
        FrameSourceTypes.Color |
        FrameSourceTypes.Depth |
        FrameSourceTypes.Body |
        FrameSourceTypes.BodyIndex);

// Subscribe to frame events
this.multiReader.MultiSourceFrameArrived +=
    FrameArrived;
```


MultiSourceFrameReader – Event handler

```
private void FrameArrived(object sender, MultiSourceFrameArrivedEventArgs e)
{
    // Acquire the color frame
    using (ColorFrame frame = e.ColorFrameReference.AcquireFrame())
    {
        ...
    }

    // Acquire the depth frame
    using (DepthFrame frame = e.DepthFrameReference.AcquireFrame())
    {
        ...
    }

    ...
}
```

Coordinate mapping



- Three coordinate systems

Name	Applies to	Dimensions	Units	Range	Origin
ColorSpacePoint	Color	2	pixels	1920x1080	Top left corner
DepthSpacePoint	Depth, Infrared, Body index	2	pixels	512x424	Top left corner
CameraSpacePoint	Body	3	meters	–	Infrared/depth camera

- Coordinate mapper provides conversions between each system
- Convert single or multiple points
- Many (but not all) methods require actual depth data

Coordinate mapping – Joint overlay

- Map the joints of a body to the color frame (e.g., to overlay joints on color image)
- We have multiple CameraSpacePoints, need ColorSpacePoints

```
int count = body.Joints.Count;
JointType[] jointTypes = body.Joints.Keys.ToArray();
CameraSpacePoint[] cameraPoints = body.Joints.Values.ToArray();
ColorSpacePoint[] colorPoints = new ColorSpacePoint[count];

this._sensor.CoordinateMapper.MapCameraPointsToColorSpace(
    cameraPoints, colorPoints);

for (int i = 0; i < count; ++i)
{
    DrawJoint(
        colorBitmap, colorPoints[i].X, colorPoint[i].Y, jointTypes[i]);
}
```


Coordinate mapping – Point cloud

- We have a frame of depth data, need CameraSpacePoints

```
int count = depthData.Length;  
CameraSpacePoint[] pointCloud = new CameraSpacePoint[count];  
this._sensor.CoordinateMapper.MapDepthFrameToCameraSpace(  
    depthData, pointCloud);
```

Coordinate mapping – Color point cloud

- We have a point cloud of CameraSpacePoints, need corresponding color values
- PointWithColor is defined by the application: a struct containing a CameraSpacePoint and a Color
- NOTE: Very slow as written; “unsafe” code would yield much better performance

```
ColorSpacePoint[] colorPoints = new ColorSpacePoint[count];  
this._sensor.CoordinateMapper.MapCameraPointsToColorSpace(  
    pointCloud, colorPoints);
```

```
PointWithColor[] colorPointCloud = new PointWithColor[count];  
for (int i = 0; i < count; ++i)  
{  
    colorPointCloud[i] = new PointWithColor  
    {  
        Position = pointCloud[i],  
        Color = GetColor(colorData, colorFrameDesc, colorPoints[i])  
    };  
}
```

CoordinateMapping – GetColor

```
private Color GetColor(
    byte[] colorData, FrameDescription frameDesc, ColorSpacePoint colorPoint)
{
    if (colorPoint.X < 0 || colorPoint.X >= frameDesc.Width ||
        colorPoint.Y < 0 || colorPoint.Y >= frameDesc.Height)
    {
        return Colors.Transparent;
    }
    int index = ((colorPoint.Y * frameDesc.Width) + colorPoint.X) * frameDesc.BytesPerPixel;
    return new Color
    {
        B = colorData[index],
        G = colorData[index + 1],
        R = colorData[index + 2],
        A = colorData[index + 3]
    };
}
```

Porting guidance

High-level changes

- Applications that use one sensor
 - May no longer need sensor chooser component
 - KinectSensor.Default "just works"
 - Can use KinectSensor.IsAvailable, if necessary
- Many fewer flags/options (near/far, standing/seated, resolution)
- Stream model → Source/reader model
- AllFramesReady → MultiSourceFrameReader

All applications – Data changes

- Body index has its own source (separate from depth)
- Color format conversion at time of frame processing
- 6 bodies fully tracked: no need to choose
- More body joints (neck, thumbs, hand tips)
- *Some joints at different positions than before*
(more anatomically correct, especially hip positions)
- Some joints renamed

New and renamed joints

- Code ported from V1 can safely ignore the new joints, and just use values 0 thru 19 of the JointType enum
- Most joints are the same in V1 and V2, except:

JointType Value	V1	V2
0	HipCenter	SpineBase
1	Spine	SpineMid
2	ShoulderCenter	Neck
20	–	SpineShoulder
21	–	HandTipLeft
22	–	ThumbLeft
23	–	HandTipRight
24	–	ThumbRight

COM applications

- New API model is much closer to .NET model
 - Significantly different API “shape” from COM V1 API
 - More interfaces, instead of monolithic INuiSensor interface
- Waitable event handles are provided by the API (not the application)

KINECT

for Windows

© 2013 Microsoft Corporation. All rights reserved. Microsoft, Windows, Windows Vista and other product names are or may be registered trademarks and/or trademarks in the U.S. and/or other countries.

The information herein is for informational purposes only and represents the current view of Microsoft Corporation as of the date of this presentation. Because Microsoft must respond to changing market conditions, it should not be interpreted to be a commitment on the part of Microsoft, and Microsoft cannot guarantee the accuracy of any information provided after the date of this presentation. MICROSOFT MAKES NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, AS TO THE INFORMATION IN THIS PRESENTATION.